

## ENGINEERING BIG DATA SYSTEMS

---



## ROAD ACCIDENTS IN THE USA

Satwik Kashyap

[kashyap.sa@northeastern.edu](mailto:kashyap.sa@northeastern.edu)

## Introduction

This is a countrywide auto collision dataset, which covers 49 states of the USA. The mishap information is gathered from February 2016 to June 2020, utilizing two APIs that give streaming traffic occurrence (or occasion) information. These APIs broadcast traffic information caught by an assortment of substances, for example, the US and state divisions of transportation, law implementation offices, traffic cameras, and traffic sensors inside the street organizations. Right now, there are about 3.5 million mishap records in this dataset. In this project, we are focused on performing Map Reduce with Java in a single standalone cluster on virtual machine

## Data Set

This dataset has been collected in real-time, using multiple Traffic APIs. Currently, it contains accident data that are collected from February 2016 to June 2020 for the Contiguous United States.

The columns we have considered for the map reduce analysis are listed below:

ID	Distance(mi)	Timezone	Precipitation(in)	Station
Source	Description	Airport_Code	Weather_Condition	Stop
TMC	Number	Weather_Timestamp	Amenity	Traffic_Calming
Severity	Street	Temperature(F)	Bump	Traffic_Signal
Start_Time	Side	Wind_Chill(F)	Crossing	Turning_Loop
End_Time	City	Humidity(%)	Give_Way	Sunrise_Sunset
Start_Lat	County	Pressure(in)	Junction	Civil_Twilight
Start_Lng	State	Visibility(mi)	No_Exit	Nautical_Twilight
End_Lat	Zipcode	Wind_Direction	Railway	Astronomical_Twilight
End_Lng	Country	Wind_Speed(mph)	Roundabout	

## HADOOP ANALYSIS

### Importing the Data into HDFS from Local

```
./hdfs dfs -copyFromLocal /home/nidhi/Desktop/Satwik_bigdata/usaccident/US_Accidents_June20.csv
/info7250_2020
```

```
nidhi@nidhi:~/usr/local/bin/hadoop-2.7.3/bin$ ./hdfs dfs -copyFromLocal /home/nidhi/Desktop/Satwik_bigdata/usaccident/US_Accidents_June20.csv /info7250_2020
nidhi@nidhi:~/usr/local/bin/hadoop-2.7.3/bin$
```

-rw-r--r--	nidhi	supergroup	1.24 GB	12/17/2020, 5:56:25 PM	1	128 MB	info7250_2020
drwxr-xr-x	nidhi	supergroup	0 B	10/15/2020, 10:11:37 PM	0	0 B	numbers
-rw-r--r--	nidhi	supergroup	8.54 KB	12/17/2020, 3:17:16 PM	1	128 MB	prod
-rw-r--r--	nidhi	supergroup	11.65 KB	12/17/2020, 3:03:49 PM	1	128 MB	product
-rw-r--r--	nidhi	supergroup	9.4 KB	12/17/2020, 2:55:45 PM	1	128 MB	products
-rw-r--r--	nidhi	supergroup	68 B	12/17/2020, 3:22:41 PM	1	128 MB	productzz
drwx-----	nidhi	supergroup	0 B	10/31/2020, 2:28:41 AM	0	0 B	tmp

### Q1 : Total Number of accidents on Left and Right side

It is a Map Reduce program to calculate the total number of accident on each side of the road

```
nidhi@nidhi:~/usr/local/bin/hadoop-2.7.3/bin$ ./hadoop jar /home/nidhi/Desktop/Satwik_bigdata/Q4.jar con.hadoop.finalProject.Q4.DriverClass /info7250_2020 /SatwikResults/Q4
20/12/17 20:46:56 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
20/12/17 20:46:56 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
20/12/17 20:46:57 INFO Input.FileInputFormat: Total input paths to process : 1
20/12/17 20:46:57 INFO mapreduce.JobSubmitter: number of splits:10
20/12/17 20:46:57 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1608248153408_0004
20/12/17 20:46:57 INFO Impl.YarnClientImpl: Submitted application application_1608248153408_0004
20/12/17 20:46:57 INFO mapreduce.Job: The url to track the job: http://nidhi:8088/proxy/application_1608248153408_0004/
20/12/17 20:46:57 INFO mapreduce.Job: Running job: job_1608248153408_0004
20/12/17 20:47:02 INFO mapreduce.Job: Job job_1608248153408_0004 running in uber mode : false
20/12/17 20:47:02 INFO mapreduce.Job: map 0% reduce 0%
20/12/17 20:47:21 INFO mapreduce.Job: map 14% reduce 0%
20/12/17 20:47:22 INFO mapreduce.Job: map 18% reduce 0%
20/12/17 20:47:24 INFO mapreduce.Job: map 33% reduce 0%
20/12/17 20:47:25 INFO mapreduce.Job: map 47% reduce 0%
20/12/17 20:47:26 INFO mapreduce.Job: map 60% reduce 0%
20/12/17 20:47:42 INFO mapreduce.Job: map 66% reduce 0%
20/12/17 20:47:43 INFO mapreduce.Job: map 68% reduce 0%
20/12/17 20:47:45 INFO mapreduce.Job: map 82% reduce 20%
20/12/17 20:47:46 INFO mapreduce.Job: map 93% reduce 20%
20/12/17 20:47:47 INFO mapreduce.Job: map 100% reduce 20%
20/12/17 20:47:48 INFO mapreduce.Job: map 100% reduce 37%
20/12/17 20:47:51 INFO mapreduce.Job: map 100% reduce 89%
20/12/17 20:47:52 INFO mapreduce.Job: map 100% reduce 100%
20/12/17 20:47:53 INFO mapreduce.Job: Job job_1608248153408_0004 completed successfully
20/12/17 20:47:53 INFO mapreduce.Job: Counters: 90
File System Counters
  FILE: Number of bytes read=41529579
  FILE: Number of bytes written=84366430
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=1327235002
  HDFS: Number of bytes written=25
  HDFS: Number of read operations=33
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
Job Counters
  Killed map tasks=1
  Launched map tasks=11
  Launched reduce tasks=1
  Data-local map tasks=11
  Total time spent by all maps in occupied slots (ms)=215247
  Total time spent by all reduces in occupied slots (ms)=24337
  Total time spent by all map tasks (ms)=215247
  Total time spent by all reduce tasks (ms)=24337
  Total vcore-millisecods taken by all map tasks=215247
  Total vcore-millisecods taken by all reduce tasks=24337
  Total megabyte-millisecods taken by all map tasks=220412928
```

```

Killed map tasks=1
Launched map tasks=11
Launched reduce tasks=1
Data-local map tasks=11
Total time spent by all maps in occupied slots (ms)=215247
Total time spent by all reduces in occupied slots (ms)=24337
Total time spent by all map tasks (ms)=215247
Total time spent by all reduce tasks (ms)=24337
Total vcore-milliseconds taken by all map tasks=215247
Total vcore-milliseconds taken by all reduce tasks=24337
Total megabyte-milliseconds taken by all map tasks=220412928
Total megabyte-milliseconds taken by all reduce tasks=24921088
Map-Reduce Framework
  Map input records=3513618
  Map output records=3513616
  Map output bytes=34502341
  Map output materialized bytes=41529633
  Input split bytes=1000
  Combine input records=0
  Combine output records=0
  Reduce input groups=2
  Reduce shuffle bytes=41529633
  Reduce input records=3513616
  Reduce output records=2
  Spilled Records=7027232
  Shuffled Maps =10
  Failed Shuffles=0
  Merged Map outputs=10
  GC time elapsed (ms)=9279
  CPU time spent (ms)=71680
  Physical memory (bytes) snapshot=2971590656
  Virtual memory (bytes) snapshot=20608212992
  Total committed heap usage (bytes)=2263351296
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=1327234602
File Output Format Counters
  Bytes Written=25
nidhi@nidhi: /usr/local/bin/hadoop-2.7.3/bin$ ./hadoop fs -cat /SatwikResults/Q4/part-r-00000 | head
Left    18.04%
Right   81.96%
nidhi@nidhi: /usr/local/bin/hadoop-2.7.3/bin$

```

## MAPPER

package com.mycompany.assignment\_4.Q1TotalNumberofaccidentsonLeftandRightside;

```

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

```

```
import java.io.IOException;
```

```
public class MapperClass extends Mapper<LongWritable, Text, Text, IntWritable> {
```

```

    Text sideOfRoad = new Text();
    IntWritable one = new IntWritable(1);

```

```
@Override
```

```
protected void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
```

```

    String []tokens = value.toString().split(",");
    if(!tokens[0].equals("ID")) {
        String side = tokens[14]; // 14th column in the data
        if(!side.equals(" ")) {

            if(side.equals("L")) {
                side = "Left";
            } else {
                side = "Right";
            }

            sideOfRoad.set(side);
            context.write(sideOfRoad, one);
        }
    }
}

```

```

    }
}
REDUCER
package com.mycompany.assignment_4.Q1TotalNumberOfaccidentsonLeftandRightside;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

import java.io.IOException;

public class ReducerClass extends Reducer<Text, IntWritable, Text, Text> {

    Text percentage = new Text();

    @Override
    protected void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException, InterruptedException {

        int count = 0;

        for(IntWritable v : values) {
            count += v.get();
        }

        int total = 3513617; // total number of records in the data

        double perc = ((double) count / total) * 100;
        percentage.set(String.format("%.2f", perc) + "%");

        context.write(key, percentage);
    }
}

```

**DRIVER:**

```

package com.mycompany.assignment_4.Q1TotalNumberOfaccidentsonLeftandRightside;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import java.io.IOException;

// Defining a Driver Class
public class DriverClass {

    public static void main(String[] args) throws IOException, InterruptedException, ClassNotFoundException {
        Configuration conf = new Configuration();

        Job job = Job.getInstance(conf, "PercentagePerSideOfRoad");

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
    }
}

```

```
// Driver Class
job.setJarByClass(DriverClass.class);

job.setInputFormatClass(TextInputFormat.class);

Path outDir = new Path(args[1]);
FileOutputFormat.setOutputPath(job, outDir);

// Mapper Class
job.setMapperClass(MapperClass.class);

job.setMapOutputKeyClass(Text.class);
job.setMapOutputValueClass(IntWritable.class);

// Reducer Class
job.setReducerClass(ReducerClass.class);

// Using it to create one reducer
job.setNumReduceTasks(1);

job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);

FileSystem fs = FileSystem.get(job.getConfiguration());
if(fs.exists(outDir)){
    fs.delete(outDir, true);
}
// Submitting the job and waiting for it to complete
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

## Q2. Weekday Analysis : Total Number of accidents

It is a Map Reduce program to calculate the total number of accident every day in a week

```
nldh@nldh: /usr/local/bin/hadoop-2.7.3/bin$ ./hadoop jar /home/nldh/Desktop/Satwik_bigdata/Q3/q3.jar con.hadoop.FinalProject.Q3.DriverClass /Info7250_2020 /SatwikResults/Q3
20/12/17 20:36:53 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
20/12/17 20:36:53 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
20/12/17 20:36:53 INFO InputFileInputFormat: Total input paths to process : 1
20/12/17 20:36:54 INFO mapreduce.JobSubmitter: number of splits:10
20/12/17 20:36:54 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1608248153408_0003
20/12/17 20:36:54 INFO Impl.VarnClientImpl: Submitted application application_1608248153408_0003
20/12/17 20:36:54 INFO mapreduce.Job: The url to track the job: http://nldh:8080/proxy/application_1608248153408_0003/
20/12/17 20:36:54 INFO mapreduce.Job: Running job: job_1608248153408_0003
20/12/17 20:36:59 INFO mapreduce.Job: Job job_1608248153408_0003 running in uber mode : false
20/12/17 20:36:59 INFO mapreduce.Job:  map 0% reduce 0%
20/12/17 20:37:16 INFO mapreduce.Job:  map 3% reduce 0%
20/12/17 20:37:17 INFO mapreduce.Job:  map 4% reduce 0%
20/12/17 20:37:19 INFO mapreduce.Job:  map 9% reduce 0%
20/12/17 20:37:20 INFO mapreduce.Job:  map 14% reduce 0%
20/12/17 20:37:22 INFO mapreduce.Job:  map 21% reduce 0%
20/12/17 20:37:23 INFO mapreduce.Job:  map 27% reduce 0%
20/12/17 20:37:26 INFO mapreduce.Job:  map 29% reduce 0%
20/12/17 20:37:27 INFO mapreduce.Job:  map 38% reduce 0%
20/12/17 20:37:29 INFO mapreduce.Job:  map 42% reduce 0%
20/12/17 20:37:30 INFO mapreduce.Job:  map 47% reduce 0%
20/12/17 20:37:31 INFO mapreduce.Job:  map 60% reduce 0%
20/12/17 20:37:49 INFO mapreduce.Job:  map 62% reduce 0%
20/12/17 20:37:51 INFO mapreduce.Job:  map 65% reduce 0%
20/12/17 20:37:52 INFO mapreduce.Job:  map 70% reduce 20%
20/12/17 20:37:54 INFO mapreduce.Job:  map 75% reduce 20%
20/12/17 20:37:55 INFO mapreduce.Job:  map 82% reduce 20%
20/12/17 20:37:56 INFO mapreduce.Job:  map 85% reduce 20%
20/12/17 20:37:57 INFO mapreduce.Job:  map 97% reduce 20%
20/12/17 20:37:58 INFO mapreduce.Job:  map 100% reduce 27%
20/12/17 20:38:01 INFO mapreduce.Job:  map 100% reduce 67%
20/12/17 20:38:04 INFO mapreduce.Job:  map 100% reduce 100%
20/12/17 20:38:05 INFO mapreduce.Job: Job job_1608248153408_0003 completed successfully
20/12/17 20:38:05 INFO mapreduce.Job: Counters: 50
  File System Counters
    FILE: Number of bytes read=49853785
    FILE: Number of bytes written=101016965
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=1327235682
    HDFS: Number of bytes written=106
    HDFS: Number of read operations=33
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Killed map tasks=2
    Launched map tasks=11
    Launched reduce tasks=1
```

```

Total time spent by all reduces in occupied slots (ms)=31635
Total time spent by all map tasks (ms)=290099
Total time spent by all reduce tasks (ms)=31635
Total vcore-millisecond taken by all map tasks=290099
Total vcore-millisecond taken by all reduce tasks=31635
Total megabyte-millisecond taken by all map tasks=297061376
Total megabyte-millisecond taken by all reduce tasks=32394240

Map-Reduce Framework
Map input records=3513618
Map output records=3513618
Map output bytes=42826543
Map output materialized bytes=49853839
Input split bytes=1000
Combine input records=0
Combine output records=0
Reduce input groups=7
Reduce shuffle bytes=49853839
Reduce input records=3513618
Reduce output records=7
Spilled Records=7027236
Shuffled Maps =10
Failed Shuffles=0
Merged Map outputs=10
GC time elapsed (ms)=11030
CPU time spent (ms)=118390
Physical memory (bytes) snapshot=3010207744
Virtual memory (bytes) snapshot=20616167424
Total committed heap usage (bytes)=2251816960

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=1327234602
File Output Format Counters
  Bytes Written=106

nidhi@nidhi: /usr/local/bin/hadoop-2.7.3/bin$ ./hadoop fs -cat /SatwikResults/Q3/part-r-00000 | head
Friday 639706
Monday 592871
Saturday 214483
Sunday 189315
Thursday 621678
Tuesday 631136
Wednesday 624429

```

```

Bytes Written=106
nidhi@nidhi: /usr/local/bin/hadoop-2.7.3/bin$ ./hadoop fs -cat /SatwikResults/Q3/part-r-00000 | head
Friday 639706
Monday 592871
Saturday 214483
Sunday 189315
Thursday 621678
Tuesday 631136
Wednesday 624429

```

## MAPPER

package com.mycompany.assignment\_4.Q2WeekdayAnalysisTotalNumberOfaccidents;

```

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

```

```
import java.io.IOException;
```

```
public class DriverClass {
```

```
    public static void main(String[] args) throws IOException, InterruptedException, ClassNotFoundException {
        Configuration conf = new Configuration();
```

```
        Job job = Job.getInstance(conf, "NumberOfAccidentsPerWeekday");
```

```
        FileInputFormat.addInputPath(job, new Path(args[0]));
```

```
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

```
        //job.getConfiguration().set("mapreduce.output.textoutputformat.recordseparator","\\t");
```

```
        // Driver Class
```

```
        job.setJarByClass(DriverClass.class);
```

```
        job.setInputFormatClass(TextInputFormat.class);
```

```
Path outDir = new Path(args[1]);
FileOutputFormat.setOutputPath(job, outDir);

// Comparator
job.setSortComparatorClass(ComparatorClass.class);

// Mapper
job.setMapperClass(MapperClass.class);

job.setMapOutputKeyClass(Text.class);
job.setMapOutputValueClass(IntWritable.class);

// Reducer
job.setReducerClass(ReducerClass.class);

// Create only 1 reducer
job.setNumReduceTasks(1);

job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);

FileSystem fs = FileSystem.get(job.getConfiguration());
if(fs.exists(outDir)){
    fs.delete(outDir, true);
}

// Submit the job, then poll for progress until the job is complete
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}
```

## REDUCER

```
package com.mycompany.assignment_4.Q2WeekdayAnalysisTotalNumberofaccidents;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

import java.io.IOException;

public class ReducerClass extends Reducer<Text, IntWritable, Text, LongWritable> {

    LongWritable sum = new LongWritable();

    @Override
    protected void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException, InterruptedException {

        long count = 0;

        for(IntWritable v : values) {
            count += v.get();
        }

        sum.set(count);

        context.write(key, sum);
    }
}
```



```
}  
}  
COMPARATOR  
package com.mycompany.assignment_4.Q2WeekdayAnalysisTotalNumberofaccidents;
```

```
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.io.WritableComparator;
```

```
import java.text.ParseException;  
import java.text.SimpleDateFormat;  
import java.util.Date;
```

```
public class ComparatorClass extends WritableComparator {
```

```
    SimpleDateFormat parser = new SimpleDateFormat("EEEE");
```

```
    protected ComparatorClass() {  
        super(Text.class, true);  
    }
```

```
    @Override
```

```
    public int compare(Object a, Object b) {  
        Text m1 = (Text) a;  
        Text m2 = (Text) b;
```

```
        Date d1 = new Date();  
        Date d2 = new Date();
```

```
        try {  
            d1 = parser.parse(m1.toString());  
            d2 = parser.parse(m2.toString());  
        } catch (ParseException e) {  
            e.printStackTrace();  
        }
```

```
        return d1.compareTo(d2);  
    }
```

```
}  
DRIVER
```

```
package com.mycompany.assignment_4.Q2WeekdayAnalysisTotalNumberofaccidents;
```

```
import org.apache.hadoop.conf.Configuration;  
import org.apache.hadoop.fs.FileSystem;  
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Job;  
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;  
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;  
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

```
import java.io.IOException;
```

```
public class DriverClass {
```

```
    public static void main(String[] args) throws IOException, InterruptedException, ClassNotFoundException {  
        Configuration conf = new Configuration();
```

```
        Job job = Job.getInstance(conf, "NumberOfAccidentsPerWeekday");
```

```
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));

//job.getConfiguration().set("mapreduce.output.textoutputformat.recordseparator","\t");

// Driver Class
job.setJarByClass(DriverClass.class);

job.setInputFormatClass(TextInputFormat.class);

Path outDir = new Path(args[1]);
FileOutputFormat.setOutputPath(job, outDir);

// Comparator
job.setSortComparatorClass(ComparatorClass.class);

// Mapper
job.setMapperClass(MapperClass.class);

job.setMapOutputKeyClass(Text.class);
job.setMapOutputValueClass(IntWritable.class);

// Reducer
job.setReducerClass(ReducerClass.class);

// Create only 1 reducer
job.setNumReduceTasks(1);

job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);

FileSystem fs = FileSystem.get(job.getConfiguration());
if(fs.exists(outDir)){
    fs.delete(outDir, true);
}

// Submit the job, then poll for progress until the job is complete
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

### Q3. Monthly Analysis : Total Number of accidents

---

It is a Map Reduce program to calculate the total number of accident in a given month

```

nidhi@nidhi:~/usr/local/bin/hadoop-2.7.3/bin$ ./hadoop jar /home/nidhi/Desktop/satwik_bigdata/satwik_info7250.jar DriverClass /info7250_2020 /satwikResults/Q1
20/12/17 18:38:58 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
20/12/17 18:38:59 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
20/12/17 18:39:00 INFO Input.FileInputFormat: Total input paths to process : 1
20/12/17 18:39:00 INFO mapreduce.JobSubmitter: number of splits:10
20/12/17 18:39:00 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1608248153408_0001
20/12/17 18:39:01 INFO InputVarlenClientImpl: Submitted application application_1608248153408_0001
20/12/17 18:39:01 INFO mapreduce.Job: The url to track the job: http://nidhi:8088/proxy/application_1608248153408_0001/
20/12/17 18:39:01 INFO mapreduce.Job: Running job: job_1608248153408_0001
20/12/17 18:39:09 INFO mapreduce.Job: Job job_1608248153408_0001 running in uber mode : false
20/12/17 18:39:09 INFO mapreduce.Job: map 0% reduce 0%
20/12/17 18:39:28 INFO mapreduce.Job: map 2% reduce 0%
20/12/17 18:39:31 INFO mapreduce.Job: map 5% reduce 0%
20/12/17 18:39:32 INFO mapreduce.Job: map 6% reduce 0%
20/12/17 18:39:34 INFO mapreduce.Job: map 11% reduce 0%
20/12/17 18:39:35 INFO mapreduce.Job: map 15% reduce 0%
20/12/17 18:39:37 INFO mapreduce.Job: map 21% reduce 0%
20/12/17 18:39:38 INFO mapreduce.Job: map 28% reduce 0%
20/12/17 18:39:40 INFO mapreduce.Job: map 29% reduce 0%
20/12/17 18:39:41 INFO mapreduce.Job: map 39% reduce 0%
20/12/17 18:39:42 INFO mapreduce.Job: map 47% reduce 0%
20/12/17 18:39:43 INFO mapreduce.Job: map 60% reduce 0%
20/12/17 18:39:57 INFO mapreduce.Job: map 62% reduce 0%
20/12/17 18:39:59 INFO mapreduce.Job: map 64% reduce 0%
20/12/17 18:40:00 INFO mapreduce.Job: map 66% reduce 0%
20/12/17 18:40:01 INFO mapreduce.Job: map 68% reduce 20%
20/12/17 18:40:02 INFO mapreduce.Job: map 73% reduce 20%
20/12/17 18:40:03 INFO mapreduce.Job: map 76% reduce 20%
20/12/17 18:40:04 INFO mapreduce.Job: map 79% reduce 20%
20/12/17 18:40:05 INFO mapreduce.Job: map 81% reduce 20%
20/12/17 18:40:06 INFO mapreduce.Job: map 89% reduce 20%
20/12/17 18:40:07 INFO mapreduce.Job: map 100% reduce 23%
20/12/17 18:40:10 INFO mapreduce.Job: map 100% reduce 63%
20/12/17 18:40:13 INFO mapreduce.Job: map 100% reduce 100%
20/12/17 18:40:14 INFO mapreduce.Job: Job job_1608248153408_0001 completed successfully
20/12/17 18:40:14 INFO mapreduce.Job: Counters: 50
  File System Counters
    FILE: Number of bytes read=46397821
    FILE: Number of bytes written=94105840
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=1327234602
    HDFS: Number of bytes written=170
    HDFS: Number of read operations=33
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Killed map tasks=1
  Map-Reduce Framework
    Map input records=3513618
    Map output records=3513618
    Map output bytes=39370579
    Map output materialized bytes=46397875
    Input split bytes=1000
    Combine input records=0
    Combine output records=0
    Reduce input groups=12
    Reduce shuffle bytes=46397875
    Reduce input records=3513618
    Reduce output records=12
    Spilled Records=7027236
    Shuffled Maps =10
    Failed Shuffles=0
    Merged Map outputs=10
    GC time elapsed (ms)=11625
    CPU time spent (ms)=118900
    Physical memory (bytes) snapshot=3051593728
    Virtual memory (bytes) snapshot=20609798144
    Total committed heap usage (bytes)=2265448448
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=1327234602
  File Output Format Counters
    Bytes Written=170
nidhi@nidhi:~/usr/local/bin/hadoop-2.7.3/bin$ ./hadoop fs -cat /satwikResults/Q1/part-r-00000 | head
April 299498
August 288930
December 299624
February 284397
January 301924
July 222968
June 310322
March 293389
May 296545
November 299056
nidhi@nidhi:~/usr/local/bin/hadoop-2.7.3/bin$

```

## MAPPER

package com.mycompany.assignment\_4.Q3MonthlyAnalysisTotalNumberOfaccidents;

```

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

```

```

import java.io.IOException;
import java.text.ParseException;
import java.text.SimpleDateFormat;
import java.util.Date;

```

```
public class MapperClass extends Mapper<LongWritable, Text, Text, IntWritable> {
```

```
Text month = new Text();
IntWritable one = new IntWritable(1);

SimpleDateFormat parser = new SimpleDateFormat("yyyy-MM-dd HH:mm:ss");
SimpleDateFormat format = new SimpleDateFormat("MMMMM"); //taking the month name

@Override
protected void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {

    String []tokens = value.toString().split(",");
    Date date = new Date();

    String timeStamp = tokens[4];

    try {
        date = parser.parse(timeStamp);
    } catch (ParseException e) {
    }

    month.set(format.format(date));

    context.write(month, one);

}
}

REDUCER
package com.mycompany.assignment_4.Q3MonthlyAnalysisTotalNumberofaccidents;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

import java.io.IOException;

public class ReducerClass extends Reducer<Text, IntWritable, Text, LongWritable> {

    LongWritable sum = new LongWritable();

    @Override
    protected void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException, InterruptedException {

        long count = 0;

        for(IntWritable v : values) {
            count += v.get();
        }

        sum.set(count);

        context.write(key, sum);
    }
}

DRIVER
package com.mycompany.assignment_4.Q3MonthlyAnalysisTotalNumberofaccidents;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
```

```
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import java.io.IOException;

public class DriverClass {

    public static void main(String[] args) throws IOException, InterruptedException, ClassNotFoundException {
        Configuration conf = new Configuration();

        Job job = Job.getInstance(conf, "NumberOfAccidentsPerMonth");

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        job.getConfiguration().set("mapreduce.output.textoutputformat.recordseparator", "\t");

        // Driver Class
        job.setJarByClass(DriverClass.class);

        job.setInputFormatClass(TextInputFormat.class);

        Path outDir = new Path(args[1]);
        FileOutputFormat.setOutputPath(job, outDir);

        // Comparator
        job.setSortComparatorClass(ComparatorClass.class);

        // Mapper
        job.setMapperClass(MapperClass.class);

        job.setMapOutputKeyClass(Text.class);
        job.setMapOutputValueClass(IntWritable.class);

        // Reducer
        job.setReducerClass(ReducerClass.class);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        // Create only 1 reducer
        job.setNumReduceTasks(1);

        FileSystem fs = FileSystem.get(job.getConfiguration());

        if(fs.exists(outDir)){
            fs.delete(outDir, true);
        }

        // Submit the job, then poll for progress until the job is complete
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```

**COMPARATOR**

```
package com.mycompany.assignment_4.Q3MonthlyAnalysisTotalNumberOfaccidents;
```

```
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.WritableComparable;
import org.apache.hadoop.io.WritableComparator;
```

```
import java.text.ParseException;
import java.text.SimpleDateFormat;
import java.util.Date;
```

```
public class ComparatorClass extends WritableComparator {
```

```
    SimpleDateFormat parser = new SimpleDateFormat("MMMMM");
```

```
    protected ComparatorClass() {
        super(Text.class, true);
    }
```

```
    @Override
```

```
    public int compare(Object a, Object b) {
        Text m1 = (Text) a;
        Text m2 = (Text) b;
```

```
        Date d1 = new Date();
        Date d2 = new Date();
```

```
        try {
            d1 = parser.parse(m1.toString());
            d2 = parser.parse(m2.toString());
        } catch (ParseException e) {
            e.printStackTrace();
        }
```

```
        return d1.compareTo(d2);
    }
```

```
}
```

**Q4. State Analysis : Total Number of accidents – top 10 prone**

It is a Map Reduce program to calculate the top 10 accident prone states

```
stahm@nidhi:~/usr/local/bin/hadoop-2.7.3/bin$ ./hadoop jar /home/nidhi/desktop/Satwik bigdata/Q6.jar com.hadoop.finalProject.Q6.DriverClass /info7250_2020 /SatwikResults/Q6
20/12/17 21:01:19 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
20/12/17 21:01:20 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool Interface and execute your application with ToolRunner to remedy this.
20/12/17 21:01:20 INFO input.FileInputFormat: Total input paths to process : 1
20/12/17 21:01:20 INFO mapreduce.JobSubmitter: number of splits:10
20/12/17 21:01:21 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1608248153408_0006
20/12/17 21:01:21 INFO impl.YarnClientImpl: Submitted application application_1608248153408_0006
20/12/17 21:01:21 INFO mapreduce.Job: The url to track the job: http://nidhi:8088/proxy/application_1608248153408_0006/
20/12/17 21:01:21 INFO mapreduce.Job: Running job: job_1608248153408_0006
20/12/17 21:01:27 INFO mapreduce.Job: Job job_1608248153408_0006 running in uber mode : false
20/12/17 21:01:27 INFO mapreduce.Job:  map 0% reduce 0%
20/12/17 21:01:46 INFO mapreduce.Job:  map 17% reduce 0%
20/12/17 21:01:47 INFO mapreduce.Job:  map 32% reduce 0%
20/12/17 21:01:49 INFO mapreduce.Job:  map 35% reduce 0%
20/12/17 21:01:50 INFO mapreduce.Job:  map 42% reduce 0%
20/12/17 21:01:52 INFO mapreduce.Job:  map 47% reduce 0%
20/12/17 21:01:53 INFO mapreduce.Job:  map 60% reduce 0%
20/12/17 21:02:10 INFO mapreduce.Job:  map 67% reduce 0%
20/12/17 21:02:11 INFO mapreduce.Job:  map 74% reduce 0%
20/12/17 21:02:12 INFO mapreduce.Job:  map 87% reduce 0%
20/12/17 21:02:13 INFO mapreduce.Job:  map 90% reduce 0%
20/12/17 21:02:14 INFO mapreduce.Job:  map 93% reduce 17%
20/12/17 21:02:15 INFO mapreduce.Job:  map 100% reduce 17%
20/12/17 21:02:17 INFO mapreduce.Job:  map 100% reduce 23%
20/12/17 21:02:20 INFO mapreduce.Job:  map 100% reduce 27%
```

```

nldht@nldht:~/usr/local/bin/hadoop-2.7.3/bin$ ./hadoop fs -tail /SatwikResults/Q6/part-r-00000 | head
California 329284
Texas 258002
Florida 258002
South Carolina 173277
North Carolina 165958
New York 160817
Pennsylvania 106787
Illinois 99692
Virginia 96075
Michigan 95983nldht@nldht:~/usr/local/bin/hadoop-2.7.3/bin$

```

#### MAPPER

```
package com.mycompany.assignment_4.Q4StateAnalysisTotalNumberOfaccidentstop10prone;
```

```
import org.apache.hadoop.io.NullWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
```

```
import java.io.IOException;
```

```
public class MapperClass extends Mapper<Object, Text, NullWritable, Text> {
```

```
    @Override
    protected void map(Object key, Text value, Context context) throws IOException, InterruptedException {
        context.write(NullWritable.get(),value);
    }
}
```

#### REDUCER

```
package com.mycompany.assignment_4.Q4StateAnalysisTotalNumberOfaccidentstop10prone;
```

```
import org.apache.hadoop.io.NullWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;
```

```
import java.io.IOException;
import java.util.Comparator;
import java.util.TreeMap;
```

```
public class ReducerClass extends Reducer<NullWritable, Text, NullWritable, Text> {
```

```
    static class DescOrder implements Comparator<Integer> {
```

```
        @Override
        public int compare(Integer o1, Integer o2) {
            return o2.compareTo(o1);
        }
    }
```

```
    private TreeMap<Integer, Text> countMap = new TreeMap<>(new DescOrder());
```

```
    @Override
    protected void reduce(NullWritable key, Iterable<Text> values, Context context) throws IOException, InterruptedException {
        for (Text value : values) {
```

```
            String[] tokens = value.toString().split("\t");
            int countOfStates = Integer.parseInt(tokens[1]);
```

```
            countMap.put(countOfStates, new Text(value));
```

```
        if (countMap.size() > 10)
            countMap.remove(countMap.lastKey());
    }

    for (Text t : countMap.values())
        context.write(NullWritable.get(), t);
    }
}

DRIVER
package com.mycompany.assignment_4.Q4StateAnalysisTotalNumberofaccidentstop10prone;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.NullWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import java.io.IOException;

public class DriverClass {

    public static void main(String[] args) throws IOException, InterruptedException, ClassNotFoundException{

        Configuration conf = new Configuration();
        // Create a new Job
        Job job = Job.getInstance(conf, "wordcount");
        job.setJarByClass(DriverClass.class);

        // Specify various job-specific parameters
        job.setJobName("myjob");

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        job.setInputFormatClass(TextInputFormat.class);

        Path outDir = new Path(args[1]);
        FileOutputFormat.setOutputPath(job, outDir);

        job.setMapOutputKeyClass(NullWritable.class);
        job.setMapOutputValueClass(Text.class);

        job.setMapperClass(MapperClass.class);
        job.setReducerClass(ReducerClass.class);

        job.setOutputKeyClass(NullWritable.class);
        job.setOutputValueClass(Text.class);

        // Create only 1 reducer
        job.setNumReduceTasks(1);

        FileSystem fs = FileSystem.get(job.getConfiguration());
```



```

if(fs.exists(outDir)){
    fs.delete(outDir, true);
}

// Submit the job, then poll for progress until the job is complete
System.exit(job.waitForCompletion(true)?0:1);
}
}

```

## Q5. Hourly Analysis : Total Number of accidents

It is a Map Reduce program to calculate the total number of accident in a given hour

```

nidhi@nidhi:~/usr/local/bin/hadoop-2.7.3/bin$ ./hadoop jar /home/nidhi/Desktop/Satwik_bigdata/Q2/BigDataFinalProject.jar com.hadoop.finalProject.Q2.DriverClass /info7250_2020 /SatwikResults/Q2
20/12/17 20:28:35 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
20/12/17 20:28:35 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
20/12/17 20:28:35 INFO InputFileInputFormat: Total input paths to process : 1
20/12/17 20:28:36 INFO mapreduce.JobSubmitter: number of splits:10
20/12/17 20:28:36 INFO mapreduce.JobSubmitter: Submitting tokens for job: Job_1608248153408_0002
20/12/17 20:28:36 INFO ImplYarnClientImpl: Submitted application application_1608248153408_0002
20/12/17 20:28:36 INFO mapreduce.Job: The url to track the job: http://nidhi:8088/proxy/application_1608248153408_0002/
20/12/17 20:28:36 INFO mapreduce.Job: Running job: job_1608248153408_0002
20/12/17 20:28:41 INFO mapreduce.Job: Job job_1608248153408_0002 running in uber mode : false
20/12/17 20:29:00 INFO mapreduce.Job: map 0% reduce 0%
20/12/17 20:29:00 INFO mapreduce.Job: map 3% reduce 0%
20/12/17 20:29:01 INFO mapreduce.Job: map 9% reduce 0%
20/12/17 20:29:03 INFO mapreduce.Job: map 12% reduce 0%
20/12/17 20:29:04 INFO mapreduce.Job: map 23% reduce 0%
20/12/17 20:29:06 INFO mapreduce.Job: map 26% reduce 0%
20/12/17 20:29:07 INFO mapreduce.Job: map 39% reduce 0%
20/12/17 20:29:08 INFO mapreduce.Job: map 46% reduce 0%
20/12/17 20:29:09 INFO mapreduce.Job: map 60% reduce 0%
20/12/17 20:29:23 INFO mapreduce.Job: map 62% reduce 0%
20/12/17 20:29:25 INFO mapreduce.Job: map 64% reduce 0%
20/12/17 20:29:26 INFO mapreduce.Job: map 70% reduce 0%
20/12/17 20:29:27 INFO mapreduce.Job: map 70% reduce 20%
20/12/17 20:29:28 INFO mapreduce.Job: map 73% reduce 20%
20/12/17 20:29:29 INFO mapreduce.Job: map 82% reduce 20%
20/12/17 20:29:31 INFO mapreduce.Job: map 88% reduce 20%
20/12/17 20:29:32 INFO mapreduce.Job: map 97% reduce 20%
20/12/17 20:29:33 INFO mapreduce.Job: map 100% reduce 30%
20/12/17 20:29:36 INFO mapreduce.Job: map 100% reduce 67%
20/12/17 20:29:38 INFO mapreduce.Job: map 100% reduce 100%
20/12/17 20:29:39 INFO mapreduce.Job: Job job_1608248153408_0002 completed successfully
20/12/17 20:29:39 INFO mapreduce.Job: Counters: 50
File System Counters
  FILE: Number of bytes read=31622568
  FILE: Number of bytes written=64550269
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=1327235602
  HDFS: Number of bytes written=230
  HDFS: Number of read operations=33
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
Job Counters
  Killed map tasks=1
  Launched map tasks=11
  Launched reduce tasks=1
  Late-local map tasks=11
  Total time spent by all maps in occupied slots (ms)=249094
  Total vcore-milliseconds taken by all map tasks=249094
  Total vcore-milliseconds taken by all reduce tasks=27439
  Total megabyte-milliseconds taken by all map tasks=255072256
  Total megabyte-milliseconds taken by all reduce tasks=28097536
Map-Reduce Framework
  Map input records=3513618
  Map output records=3513618
  Map output bytes=24595326
  Map output materialized bytes=31622622
  Input split bytes=1000
  Combine input records=0
  Combine output records=0
  Reduce input groups=24
  Reduce shuffle bytes=31622622
  Reduce input records=3513618
  Reduce output records=24
  Spilled Records=7027236
  Shuffled Maps =10
  Failed Shuffles=0
  Merged Map outputs=10
  GC time elapsed (ms)=10512
  CPU time spent (ms)=100180
  Physical memory (bytes) snapshot=2968145920
  Virtual memory (bytes) snapshot=20614549504
  Total committed heap usage (bytes)=2197815296
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=1327234602
File Output Format Counters
  Bytes Written=230
nidhi@nidhi:~/usr/local/bin/hadoop-2.7.3/bin$ ./hadoop fs -cat /SatwikResults/Q2/part-r-00000 | head
00 29109
01 22401
02 23120
03 22702
04 64706
05 99244
06 198679
07 316352
08 326257
09 202664

```

## MAPPER

```
package com.mycompany.assignment_4.Q5HourlyAnalysisTotalNumberofaccidents;
```

```
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.LongWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Mapper;
```

```
import java.io.IOException;  
import java.text.ParseException;  
import java.text.SimpleDateFormat;  
import java.util.Date;
```

```
public class MapperClass extends Mapper<LongWritable, Text, Text, IntWritable> {
```

```
    Text hour = new Text();  
    IntWritable one = new IntWritable(1);
```

```
    SimpleDateFormat parser = new SimpleDateFormat("yyyy-MM-dd HH:mm:ss");  
    SimpleDateFormat format = new SimpleDateFormat("HH");
```

```
    @Override
```

```
    protected void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
```

```
        String []tokens = value.toString().split(",");  
        Date date = new Date();
```

```
        if(!tokens[0].equals("ID")) {  
            String timeStamp = tokens[4];
```

```
            try {  
                date = parser.parse(timeStamp);  
            } catch (ParseException e) {  
                e.printStackTrace();  
            }  
        }
```

```
        hour.set(format.format(date));
```

```
        context.write(hour, one);
```

```
    }  
}
```

## REDUCER

```
package com.mycompany.assignment_4.Q5HourlyAnalysisTotalNumberofaccidents;
```

```
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.LongWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Reducer;
```

```
import java.io.IOException;
```

```
public class ReducerClass extends Reducer<Text, IntWritable, Text, LongWritable> {
```

```
    LongWritable sum = new LongWritable();
```

```

@Override
protected void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException, InterruptedException {

    long count = 0;

    for(IntWritable v : values) {
        count += v.get();
    }

    sum.set(count);

    context.write(key, sum);
}
}

```

#### COMPARATOR

```
package com.mycompany.assignment_4.Q5HourlyAnalysisTotalNumberofaccidents;
```

```
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.WritableComparator;
```

```
import java.text.ParseException;
import java.text.SimpleDateFormat;
import java.util.Date;
```

```
public class ComparatorClass extends WritableComparator {
```

```
    SimpleDateFormat parser = new SimpleDateFormat("HH");
```

```
    protected ComparatorClass() {
        super(Text.class, true);
    }

```

```

@Override
public int compare(Object a, Object b) {
    Text m1 = (Text) a;
    Text m2 = (Text) b;

    Date d1 = new Date();
    Date d2 = new Date();

    try {
        d1 = parser.parse(m1.toString());
        d2 = parser.parse(m2.toString());
    } catch (ParseException e) {
        e.printStackTrace();
    }

    return d1.compareTo(d2);
}
}

```

#### DRIVER

```
package com.mycompany.assignment_4.Q5HourlyAnalysisTotalNumberofaccidents;
```

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
```

```
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import java.io.IOException;

public class DriverClass {

    public static void main(String[] args) throws IOException, InterruptedException, ClassNotFoundException {
        Configuration conf = new Configuration();

        Job job = Job.getInstance(conf, "NumberOfAccidentsPerHourOfDay");

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        job.getConfiguration().set("mapreduce.output.textoutputformat.recordseparator", "\t");

        // Driver Class
        job.setJarByClass(DriverClass.class);

        job.setInputFormatClass(TextInputFormat.class);

        Path outDir = new Path(args[1]);
        FileOutputFormat.setOutputPath(job, outDir);

        // Comparator
        job.setSortComparatorClass(ComparatorClass.class);

        // Mapper
        job.setMapperClass(MapperClass.class);

        job.setMapOutputKeyClass(Text.class);
        job.setMapOutputValueClass(IntWritable.class);

        // Reducer
        job.setReducerClass(ReducerClass.class);

        // Create only 1 reducer
        job.setNumReduceTasks(1);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        FileSystem fs = FileSystem.get(job.getConfiguration());
        if(fs.exists(outDir)){
            fs.delete(outDir, true);
        }

        // Submit the job, then poll for progress until the job is complete
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```

## Q6. State Analysis : Total Number of accidents

It is a Map Reduce program to calculate the total number of accident per state in the USA

```

nidhi@nidhi:~/usr/local/bin/hadoop-2.7.3/bin$ ./hadoop jar /home/nidhi/Desktop/Satwik_bgdata/Q5.jar com.hadoop.finalProject.Q5.DriverClass /Info7250_2020 /SatwikResults/Q5
20/12/17 20:53:09 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
20/12/17 20:53:09 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
20/12/17 20:53:10 INFO Input.FileInputFormat: Total input paths to process : 1
20/12/17 20:53:10 INFO mapreduce.JobSubmitter: number of splits:10
20/12/17 20:53:11 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1608248153408_0005
20/12/17 20:53:11 INFO Inpl.VarnClientImpl: Submitted application application_1608248153408_0005
20/12/17 20:53:11 INFO mapreduce.Job: The url to track the job: http://nidhi:8088/proxy/application_1608248153408_0005/
20/12/17 20:53:11 INFO mapreduce.Job: Running job: job_1608248153408_0005
20/12/17 20:53:18 INFO mapreduce.Job: Job job_1608248153408_0005 running in uber mode : false
20/12/17 20:53:18 INFO mapreduce.Job: map 0% reduce 0%
20/12/17 20:53:36 INFO mapreduce.Job: map 8% reduce 0%
20/12/17 20:53:37 INFO mapreduce.Job: map 23% reduce 0%
20/12/17 20:53:39 INFO mapreduce.Job: map 26% reduce 0%
20/12/17 20:53:40 INFO mapreduce.Job: map 37% reduce 0%
20/12/17 20:53:41 INFO mapreduce.Job: map 48% reduce 0%
20/12/17 20:53:42 INFO mapreduce.Job: map 68% reduce 0%
20/12/17 20:53:59 INFO mapreduce.Job: map 74% reduce 0%
20/12/17 20:54:02 INFO mapreduce.Job: map 87% reduce 20%
20/12/17 20:54:03 INFO mapreduce.Job: map 100% reduce 20%
20/12/17 20:54:06 INFO mapreduce.Job: map 100% reduce 67%
20/12/17 20:54:08 INFO mapreduce.Job: map 100% reduce 100%
20/12/17 20:54:08 INFO mapreduce.Job: Job job_1608248153408_0005 completed successfully
20/12/17 20:54:08 INFO mapreduce.Job: Counters: 50
  File System Counters
    FILE: Number of bytes read=55504286
    FILE: Number of bytes written=112315866
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=1327235602
    HDFS: Number of bytes written=761
    HDFS: Number of read operations=33
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Killed map tasks=1
    Launched map tasks=11
    Launched reduce tasks=1
    Data-local map tasks=11
    Total time spent by all maps in occupied slots (ms)=218410
    Total time spent by all reduces in occupied slots (ms)=22500
    Total time spent by all map tasks (ms)=218410
    Total time spent by all reduce tasks (ms)=22500
    Total vcore-milliseconds taken by all map tasks=218410
    Total vcore-milliseconds taken by all reduce tasks=22500
    Total megabyte-milliseconds taken by all map tasks=223651840
    Total megabyte-milliseconds taken by all reduce tasks=23040000

```

```

nidhi@nidhi:~/usr/local/bin/hadoop-2.7.3/bin$ ./hadoop fs -cat /SatwikResults/Q5/part-r-00000 | head
1
Alabama 44625
Arizona 78584
Arkansas 2012
California 816825
Colorado 49731
Connecticut 25901
Delaware 5739
District Of Columbia 4820
Florida 258002
nidhi@nidhi:~/usr/local/bin/hadoop-2.7.3/bin$

```

### MAPPER

package com.mycompany.assignment\_4.Q6StateAnalysisTotalNumberOfaccidents;

```

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

```

```
import java.io.IOException;
```

```
public class MapperClass extends Mapper<LongWritable, Text, Text, IntWritable> {
```

```
    Text state = new Text();
```

```
    IntWritable one = new IntWritable(1);
```

```
    StateMap statesMap = new StateMap();
```

```
@Override
```

```
protected void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
```

```
    String []tokens = value.toString().split(",");
```

```
    String stateString = "";
```

```
        if(!tokens[0].equals("ID")) {  
            stateString = statesMap.getStateFromAbbr(tokens[17]);  
        }  
  
        state.set(stateString);  
  
        context.write(state, one);  
    }  
}
```

#### REDUCER

```
package com.mycompany.assignment_4.Q6StateAnalysisTotalNumberofaccidents;
```

```
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.LongWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Reducer;
```

```
import java.io.IOException;
```

```
public class ReducerClass extends Reducer<Text, IntWritable, Text, LongWritable> {
```

```
    LongWritable sum = new LongWritable();
```

```
    @Override
```

```
    protected void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException, InterruptedException {
```

```
        long count = 0;
```

```
        for(IntWritable v : values) {  
            count += v.get();  
        }
```

```
        sum.set(count);
```

```
        context.write(key, sum);  
    }
```

```
}
```

#### DRIVER

```
package com.mycompany.assignment_4.Q6StateAnalysisTotalNumberofaccidents;
```

```
import org.apache.hadoop.conf.Configuration;  
import org.apache.hadoop.fs.FileSystem;  
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Job;  
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;  
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;  
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

```
import java.io.IOException;
```

```
public class DriverClass {
```

```

public static void main(String[] args) throws IOException, InterruptedException, ClassNotFoundException {
    Configuration conf = new Configuration();

    Job job = Job.getInstance(conf, "NumberOfAccidentsPerState");

    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    //job.getConfiguration().set("mapreduce.output.textoutputformat.recordseparator", "\t");

    // Driver Class
    job.setJarByClass(DriverClass.class);

    job.setInputFormatClass(TextInputFormat.class);

    Path outDir = new Path(args[1]);
    FileOutputFormat.setOutputPath(job, outDir);

    // Mapper
    job.setMapperClass(MapperClass.class);

    job.setMapOutputKeyClass(Text.class);
    job.setMapOutputValueClass(IntWritable.class);

    // Reducer
    job.setReducerClass(ReducerClass.class);
    job.setNumReduceTasks(1);

    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);

    FileSystem fs = FileSystem.get(job.getConfiguration());
    if(fs.exists(outDir)){
        fs.delete(outDir, true);
    }

    // Submit the job, then poll for progress until the job is complete
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

#### STATEMAP

```

package com.mycompany.assignment_4.Q6StateAnalysisTotalNumberofaccidents;

import java.util.HashMap;
import java.util.Map;

public class StateMap {

    public static final Map<String, String> STATE_MAP;

    static {
        STATE_MAP = new HashMap<String, String>();
        STATE_MAP.put("AL", "Alabama");
        STATE_MAP.put("AK", "Alaska");
        STATE_MAP.put("AB", "Alberta");
        STATE_MAP.put("AZ", "Arizona");
    }
}

```

```
STATE_MAP.put("AR", "Arkansas");
STATE_MAP.put("BC", "British Columbia");
STATE_MAP.put("CA", "California");
STATE_MAP.put("CO", "Colorado");
STATE_MAP.put("CT", "Connecticut");
STATE_MAP.put("DE", "Delaware");
STATE_MAP.put("DC", "District Of Columbia");
STATE_MAP.put("FL", "Florida");
STATE_MAP.put("GA", "Georgia");
STATE_MAP.put("GU", "Guam");
STATE_MAP.put("HI", "Hawaii");
STATE_MAP.put("ID", "Idaho");
STATE_MAP.put("IL", "Illinois");
STATE_MAP.put("IN", "Indiana");
STATE_MAP.put("IA", "Iowa");
STATE_MAP.put("KS", "Kansas");
STATE_MAP.put("KY", "Kentucky");
STATE_MAP.put("LA", "Louisiana");
STATE_MAP.put("ME", "Maine");
STATE_MAP.put("MB", "Manitoba");
STATE_MAP.put("MD", "Maryland");
STATE_MAP.put("MA", "Massachusetts");
STATE_MAP.put("MI", "Michigan");
STATE_MAP.put("MN", "Minnesota");
STATE_MAP.put("MS", "Mississippi");
STATE_MAP.put("MO", "Missouri");
STATE_MAP.put("MT", "Montana");
STATE_MAP.put("NE", "Nebraska");
STATE_MAP.put("NV", "Nevada");
STATE_MAP.put("NB", "New Brunswick");
STATE_MAP.put("NH", "New Hampshire");
STATE_MAP.put("NJ", "New Jersey");
STATE_MAP.put("NM", "New Mexico");
STATE_MAP.put("NY", "New York");
STATE_MAP.put("NF", "Newfoundland");
STATE_MAP.put("NC", "North Carolina");
STATE_MAP.put("ND", "North Dakota");
STATE_MAP.put("NT", "Northwest Territories");
STATE_MAP.put("NS", "Nova Scotia");
STATE_MAP.put("NU", "Nunavut");
STATE_MAP.put("OH", "Ohio");
STATE_MAP.put("OK", "Oklahoma");
STATE_MAP.put("ON", "Ontario");
STATE_MAP.put("OR", "Oregon");
STATE_MAP.put("PA", "Pennsylvania");
STATE_MAP.put("PE", "Prince Edward Island");
STATE_MAP.put("PR", "Puerto Rico");
STATE_MAP.put("QC", "Quebec");
STATE_MAP.put("RI", "Rhode Island");
STATE_MAP.put("SK", "Saskatchewan");
STATE_MAP.put("SC", "South Carolina");
STATE_MAP.put("SD", "South Dakota");
STATE_MAP.put("TN", "Tennessee");
STATE_MAP.put("TX", "Texas");
STATE_MAP.put("UT", "Utah");
STATE_MAP.put("VT", "Vermont");
STATE_MAP.put("VI", "Virgin Islands");
STATE_MAP.put("VA", "Virginia");
STATE_MAP.put("WA", "Washington");
STATE_MAP.put("WV", "West Virginia");
```



```

STATE_MAP.put("WI", "Wisconsin");
STATE_MAP.put("WY", "Wyoming");
STATE_MAP.put("YT", "Yukon Territory");
}

public String getStateFromAbbr(String abbr) {
    return STATE_MAP.get(abbr);
}
}

```

## Q7. Temperature Analysis

This MapReduce Numerical Summarization job performs analysis to determine the minimum, maximum and average temperature (F) per severity (range from 1 to 4 ) of accident.

```

nidhi@nidhi: /usr/local/bin/hadoop-2.7.3/bin$ ./hadoop jar /home/nidhi/Desktop/Satwik_bigdata/Q7.jar com.hadoop.finalProject.Q8.DriverClass /info7250_2020 /SatwikResults/Q7
20/12/17 21:18:07 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
20/12/17 21:18:08 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
20/12/17 21:18:08 INFO input.FileInputFormat: Total input paths to process : 1
20/12/17 21:18:08 INFO mapreduce.JobSubmitter: number of splits:10
20/12/17 21:18:09 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1608248153408_0007
20/12/17 21:18:09 INFO impl.VarnClientImpl: Submitted application application_1608248153408_0007
20/12/17 21:18:09 INFO mapreduce.Job: The url to track the job: http://nidhi:8088/proxy/application_1608248153408_0007/
20/12/17 21:18:09 INFO mapreduce.Job: Running job: job_1608248153408_0007
20/12/17 21:18:16 INFO mapreduce.Job: Job job_1608248153408_0007 running in uber mode : false
20/12/17 21:18:16 INFO mapreduce.Job: map 0% reduce 0%
20/12/17 21:18:41 INFO mapreduce.Job: map 6% reduce 0%
20/12/17 21:18:44 INFO mapreduce.Job: map 13% reduce 0%
20/12/17 21:18:45 INFO mapreduce.Job: map 15% reduce 0%
20/12/17 21:18:47 INFO mapreduce.Job: map 26% reduce 0%
20/12/17 21:18:48 INFO mapreduce.Job: map 30% reduce 0%
20/12/17 21:18:50 INFO mapreduce.Job: map 32% reduce 0%
20/12/17 21:18:51 INFO mapreduce.Job: map 47% reduce 0%
20/12/17 21:18:52 INFO mapreduce.Job: map 60% reduce 33%
20/12/17 21:19:11 INFO mapreduce.Job: map 61% reduce 0%
20/12/17 21:19:15 INFO mapreduce.Job: map 66% reduce 0%
20/12/17 21:19:16 INFO mapreduce.Job: map 68% reduce 20%
20/12/17 21:19:18 INFO mapreduce.Job: map 78% reduce 20%
20/12/17 21:19:19 INFO mapreduce.Job: map 85% reduce 20%
20/12/17 21:19:20 INFO mapreduce.Job: map 95% reduce 20%
20/12/17 21:19:21 INFO mapreduce.Job: map 100% reduce 20%
20/12/17 21:19:22 INFO mapreduce.Job: map 100% reduce 33%
20/12/17 21:19:25 INFO mapreduce.Job: map 100% reduce 82%
20/12/17 21:19:28 INFO mapreduce.Job: map 100% reduce 100%
20/12/17 21:19:29 INFO mapreduce.Job: Job job_1608248153408_0007 completed successfully
20/12/17 21:19:29 INFO mapreduce.Job: Counters: 50
File System Counters
  FILE: Number of bytes read=131019642
  FILE: Number of bytes written=263346880
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=1327235602
  HDFS: Number of bytes written=374
  HDFS: Number of read operations=33
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
Job Counters
  Killed map tasks=1
  Launched map tasks=11
  Launched reduce tasks=1
  Data-local map tasks=11
  Total time spent by all maps in occupied slots (ms)=320075
  Total time spent by all reduces in occupied slots (ms)=33816

```

```

HDFS: Number of large read operations=0
HDFS: Number of write operations=2
Job Counters
  Killed map tasks=1
  Launched map tasks=11
  Launched reduce tasks=1
  Data-local map tasks=11
  Total time spent by all maps in occupied slots (ms)=320075
  Total time spent by all reduces in occupied slots (ms)=33816
  Total time spent by all map tasks (ms)=320075
  Total time spent by all reduce tasks (ms)=33816
  Total vcore-milliseconds taken by all map tasks=320075
  Total vcore-milliseconds taken by all reduce tasks=33816
  Total megabyte-milliseconds taken by all map tasks=327756800
  Total megabyte-milliseconds taken by all reduce tasks=34627584
Map-Reduce Framework
  Map input records=3513618
  Map output records=3447885
  Map output bytes=124123860
  Map output materialized bytes=131019690
  Input split bytes=1000
  Combine input records=0
  Combine output records=0
  Reduce input groups=4
  Reduce shuffle bytes=131019690
  Reduce input records=3447885
  Reduce output records=4
  Spilled Records=6895770
  Shuffled Maps=10
  Failed Shuffles=0
  Merged Map outputs=10
  GC time elapsed (ms)=12984
  CPU time spent (ms)=105260
  Physical memory (bytes) snapshot=3055820800
  Virtual memory (bytes) snapshot=20606027840
  Total committed heap usage (bytes)=2263351296
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=1327234602
File Output Format Counters
  Bytes Written=374

```

```
nldht@nldht:/usr/local/bin/hadoop-2.7.3/bin$ ./hadoop fs -tail /SatwikResults/Q7/part-r-00000 | head
1 Max Temperature = 111.0 | Min Temperature = 3.9 | Average Temperature = 70.74177685233381
2 Max Temperature = 170.6 | Min Temperature = -89.0 | Average Temperature = 61.99493106361987
3 Max Temperature = 167.0 | Min Temperature = -89.0 | Average Temperature = 61.85957242491674
4 Max Temperature = 117.0 | Min Temperature = -40.0 | Average Temperature = 59.02189844306117
nldht@nldht:/usr/local/bin/hadoop-2.7.3/bin$
```

## MAPPER

```
package com.mycompany.assignment_4.Q7TemperatureAnalysis;
```

```
import org.apache.hadoop.io.IntWritable;
```

```
import org.apache.hadoop.io.Text;
```

```
import org.apache.hadoop.mapreduce.Mapper;
```

```
import java.io.IOException;
```

```
public class MapperClass extends Mapper<Object, Text, IntWritable, TempWritable> {
```

```
    IntWritable sev = new IntWritable();
```

```
    TempWritable tempWritable = new TempWritable();
```

```
    @Override
```

```
    protected void map(Object key, Text value, Context context) throws IOException, InterruptedException {
```

```
        String[] tokens = value.toString().split(",");
```

```
        if (!tokens[0].equals("ID")) {
```

```
            if(tokens[3].isEmpty() || tokens[23].isEmpty() || tokens[3].equals(" ") || tokens[26].equals(" ")) {
```

```
                return;
```

```
            }
```

```
            int severity = Integer.parseInt(tokens[3]);
```

```
            double temperature = Double.parseDouble(tokens[23]);
```

```
            tempWritable.setAverageTemp(temperature);
```

```
            tempWritable.setMaxTemp(temperature);
```

```
            tempWritable.setMinTemp(temperature);
```

```
            tempWritable.setCount(1);
```

```
            sev.set(severity);
```

```
            context.write(sev, tempWritable);
```

```
    }  
  }  
}
```

## REDUCER

```
package com.mycompany.assignment_4.Q7TemperatureAnalysis;  
  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.mapreduce.Reducer;  
  
import java.io.IOException;  
  
public class ReducerClass extends Reducer<IntWritable, TempWritable, IntWritable, TempWritable> {  
  
    TempWritable newTuple = new TempWritable();  
  
    @Override  
    protected void reduce(IntWritable key, Iterable<TempWritable> values, Context context) throws IOException,  
        InterruptedException {  
  
        double maxTemp = Double.MIN_VALUE;  
        double minTemp = Double.MAX_VALUE;  
        long count = 0;  
        double sumOfTemps = 0.0;  
  
        for (TempWritable tuple : values) {  
  
            sumOfTemps += tuple.getAverageTemp() * tuple.getCount();  
            count += tuple.getCount();  
  
            if(tuple.getMinTemp() < minTemp) {  
                minTemp = tuple.getMinTemp();  
            }  
  
            if(tuple.getMaxTemp() > maxTemp) {  
                maxTemp = tuple.getMaxTemp();  
            }  
        }  
  
        newTuple.setAverageTemp(sumOfTemps / count);  
        newTuple.setMaxTemp(maxTemp);  
        newTuple.setMinTemp(minTemp);  
        newTuple.setCount(count);  
    }  
}
```

```
        context.write(key, newTuple);
    }
}
```

### **TEMPWRITABLE**

```
package com.mycompany.assignment_4.Q7TemperatureAnalysis;
```

```
import org.apache.hadoop.io.Writable;
```

```
import java.io.DataInput;
```

```
import java.io.DataOutput;
```

```
import java.io.IOException;
```

```
public class TempWritable implements Writable {
```

```
    private double maxTemp;
```

```
    private double averageTemp;
```

```
    private double minTemp;
```

```
    public double getMinTemp() {
```

```
        return minTemp;
```

```
    }
```

```
    public void setMinTemp(double minTemp) {
```

```
        this.minTemp = minTemp;
```

```
    }
```

```
    private long count;
```

```
    public long getCount() {
```

```
        return count;
```

```
    }
```

```
    public void setCount(long count) {
```

```
        this.count = count;
```

```
    }
```

```
    public double getMaxTemp() {
```

```
        return maxTemp;
```

```
    }
```

```
public void setMaxTemp(double maxTemp) {  
    this.maxTemp = maxTemp;  
}
```

```
public double getAverageTemp() {  
    return averageTemp;  
}
```

```
public void setAverageTemp(double averageTemp) {  
    this.averageTemp = averageTemp;  
}
```

@Override

```
public void write(DataOutput dataOutput) throws IOException {  
    dataOutput.writeDouble(averageTemp);  
    dataOutput.writeDouble(maxTemp);  
    dataOutput.writeDouble(minTemp);  
    dataOutput.writeLong(count);  
}
```

@Override

```
public void readFields(DataInput dataInput) throws IOException {  
    averageTemp = dataInput.readDouble();  
    maxTemp = dataInput.readDouble();  
    minTemp = dataInput.readDouble();  
    count = dataInput.readLong();  
}
```

@Override

```
public String toString() {  
    return "Max Temperature = " + maxTemp +  
        " | Min Temperature = " + minTemp +  
        " | Average Temperature = " + averageTemp;  
}  
}
```

## **DRIVER**

```
package com.mycompany.assignment_4.Q7TemperatureAnalysis;
```

```
import org.apache.hadoop.conf.Configuration;  
import org.apache.hadoop.fs.FileSystem;  
import org.apache.hadoop.fs.Path;
```

```
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import java.io.IOException;

public class DriverClass {

    public static void main(String[] args) throws IOException, InterruptedException, ClassNotFoundException {
        Configuration conf = new Configuration();

        Job job = Job.getInstance(conf, "MaxAvePressurePerSeverity");

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        // Driver Class
        job.setJarByClass(DriverClass.class);

        job.setInputFormatClass(TextInputFormat.class);

        Path outDir = new Path(args[1]);
        FileOutputFormat.setOutputPath(job, outDir);

        // Mapper
        job.setMapperClass(MapperClass.class);

        job.setMapOutputKeyClass(IntWritable.class);
        job.setMapOutputValueClass(TempWritable.class);

        // Reducer
        job.setReducerClass(ReducerClass.class);

        // Create only 1 reducer
        job.setNumReduceTasks(1);

        job.setOutputKeyClass(IntWritable.class);
        job.setOutputValueClass(TempWritable.class);
```

```

FileSystem fs = FileSystem.get(job.getConfiguration());
if(fs.exists(outDir)){
    fs.delete(outDir, true);
}

// Submit the job, then poll for progress until the job is complete
System.exit(job.waitForCompletion(true) ? 0 : 1);
}

}

```

## Q8. Year and State Analysis : Total Number of accidents

This MapReduce Numerical Summarization job performs the number of accidents per state per job partitioned per year

```

nidhi@nidhi:~/usr/local/bin/hadoop-2.7.3/bin$ ./hadoop jar /home/nidhi/Desktop/Satwik_bigdata/Q8.jar com.hadoop.finalProject.Q8.DriverClass /info7250_2020 /SatwikResults/Q8
20/12/17 21:25:42 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
20/12/17 21:25:42 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
20/12/17 21:25:44 INFO Input.FileInputFormat: Total input paths to process : 1
20/12/17 21:25:44 INFO mapreduce.JobSubmitter: number of splits:10
20/12/17 21:25:45 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1608248153408_0008
20/12/17 21:25:46 INFO Impl.VarnClientImpl: Submitted application application_1608248153408_0008
20/12/17 21:25:46 INFO mapreduce.Job: The url to track the job: http://nidhi:8088/proxy/application_1608248153408_0008/
20/12/17 21:25:46 INFO mapreduce.Job: Running job: job_1608248153408_0008
20/12/17 21:25:54 INFO mapreduce.Job: Job job_1608248153408_0008 running in uber mode : false
20/12/17 21:25:54 INFO mapreduce.Job:  map 0% reduce 0%
20/12/17 21:26:12 INFO mapreduce.Job:  map 1% reduce 0%
20/12/17 21:26:13 INFO mapreduce.Job:  map 6% reduce 0%
20/12/17 21:26:15 INFO mapreduce.Job:  map 8% reduce 0%
20/12/17 21:26:16 INFO mapreduce.Job:  map 17% reduce 0%
20/12/17 21:26:19 INFO mapreduce.Job:  map 27% reduce 0%
20/12/17 21:26:20 INFO mapreduce.Job:  map 29% reduce 0%
20/12/17 21:26:22 INFO mapreduce.Job:  map 34% reduce 0%
20/12/17 21:26:23 INFO mapreduce.Job:  map 39% reduce 0%
20/12/17 21:26:25 INFO mapreduce.Job:  map 43% reduce 0%
20/12/17 21:26:26 INFO mapreduce.Job:  map 53% reduce 0%
20/12/17 21:26:27 INFO mapreduce.Job:  map 60% reduce 0%
20/12/17 21:26:27 INFO mapreduce.Job:  map 61% reduce 0%
20/12/17 21:26:48 INFO mapreduce.Job:  map 62% reduce 0%
20/12/17 21:26:49 INFO mapreduce.Job:  map 63% reduce 0%
20/12/17 21:26:50 INFO mapreduce.Job:  map 65% reduce 0%
20/12/17 21:26:51 INFO mapreduce.Job:  map 67% reduce 8%
20/12/17 21:26:52 INFO mapreduce.Job:  map 71% reduce 8%
20/12/17 21:26:54 INFO mapreduce.Job:  map 78% reduce 8%
20/12/17 21:26:55 INFO mapreduce.Job:  map 84% reduce 8%
20/12/17 21:26:57 INFO mapreduce.Job:  map 85% reduce 8%
20/12/17 21:26:58 INFO mapreduce.Job:  map 93% reduce 8%
20/12/17 21:26:59 INFO mapreduce.Job:  map 100% reduce 8%
20/12/17 21:27:00 INFO mapreduce.Job:  map 100% reduce 11%
20/12/17 21:27:01 INFO mapreduce.Job:  map 100% reduce 20%
20/12/17 21:27:03 INFO mapreduce.Job:  map 100% reduce 27%
20/12/17 21:27:04 INFO mapreduce.Job:  map 100% reduce 34%
20/12/17 21:27:06 INFO mapreduce.Job:  map 100% reduce 40%
20/12/17 21:27:15 INFO mapreduce.Job:  map 100% reduce 100%
20/12/17 21:27:17 INFO mapreduce.Job: Job job_1608248153408_0008 completed successfully
20/12/17 21:27:18 INFO mapreduce.Job: Counters: 51
File System Counters
  FILE: Number of bytes read=56217902
  FILE: Number of bytes written=114231089
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=1327235002
  HDFS: Number of bytes written=8535

```

```

nldhi@nldhi: /usr/local/bin/hadoop-2.7.3/bin
HDFS: Number of write operations=10
Job Counters
  Killed map tasks=1
  Killed reduce tasks=1
  Launched map tasks=10
  Launched reduce tasks=6
  Data-local map tasks=10
  Total time spent by all maps in occupied slots (ms)=310341
  Total time spent by all reduces in occupied slots (ms)=126288
  Total time spent by all map tasks (ms)=310341
  Total time spent by all reduce tasks (ms)=126288
  Total vcore-milliseconds taken by all map tasks=310341
  Total vcore-milliseconds taken by all reduce tasks=126288
  Total megabyte-milliseconds taken by all map tasks=317789184
  Total megabyte-milliseconds taken by all reduce tasks=129318912
Map-Reduce Framework
  Map input records=3513618
  Map output records=3513617
  Map output bytes=49198638
  Map output materialized bytes=56218172
  Input split bytes=1000
  Combine input records=0
  Combine output records=0
  Reduce input groups=245
  Reduce shuffle bytes=56218172
  Reduce input records=3513617
  Reduce output records=245
  Spilled Records=7027234
  Shuffled Maps =50
  Failed Shuffles=0
  Merged Map outputs=50
  GC time elapsed (ms)=14530
  CPU time spent (ms)=161800
  Physical memory (bytes) snapshot=3740712960
  Virtual memory (bytes) snapshot=26137754624
  Total committed heap usage (bytes)=2701131776
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  ID_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=1327234602
File Output Format Counters
  Bytes Written=8535
nldhi@nldhi: /usr/local/bin/hadoop-2.7.3/bin$

```

```

nldhi@nldhi: /usr/local/bin/hadoop-2.7.3/bin$ ./hadoop fs -cat /SatwikResults/Q8/part-r-000000 | head
Alabama - in Year - 2020 :      8251
Arkansas - in Year - 2020 :      263
Arizona - in Year - 2020 :     16257
California - in Year - 2020 :   153560
Colorado - in Year - 2020 :     9604
Connecticut - in Year - 2020 :   3097
District Of Columbia - in Year - 2020 :      1168
Delaware - in Year - 2020 :     1305
Florida - in Year - 2020 :     34251
Georgia - in Year - 2020 :     9990
nldhi@nldhi: /usr/local/bin/hadoop-2.7.3/bin$

```

```

nldhi@nldhi: /usr/local/bin/hadoop-2.7.3/bin$ ./hadoop fs -tail /SatwikResults/Q8/part-r-000000 | head
020 : 178
Michigan - in Year - 2020 :     7294
Minnesota - in Year - 2020 :    19131
Missouri - in Year - 2020 :     4629
Mississippi - in Year - 2020 :    625
Montana - in Year - 2020 :       8
North Carolina - in Year - 2020 :   23494
North Dakota - in Year - 2020 :     1
Nebraska - in Year - 2020 :     1464
New Hampshire - in Year - 2020 :    920
nldhi@nldhi: /usr/local/bin/hadoop-2.7.3/bin$

```

## CompositeKeyComparator

```
package com.mycompany.assignment_4.Q8YearandStateAnalysisTotalNumberofaccidents;
```

```
import org.apache.hadoop.io.WritableComparable;
```

```
import org.apache.hadoop.io.WritableComparator;
```

```
public class CompositeKeyComparator extends WritableComparator {
```

```
    protected CompositeKeyComparator() {
```

```
        super(CompositeKeyWritable.class, true);
```

```
    }
```



```
@Override
public int compare(WritableComparable a, WritableComparable b) {

    CompositeKeyWritable w1 = (CompositeKeyWritable) a;
    CompositeKeyWritable w2 = (CompositeKeyWritable) b;

    int result = w1.getYear().compareTo(w2.getYear());

    if(result == 0) {
        return w1.getState().compareTo(w2.getState());
    }

    return result;
}

}

CompositeKeyWritable
package com.mycompany.assignment_4.Q8YearandStateAnalysisTotalNumberofaccidents;

import org.apache.hadoop.io.WritableComparable;

import java.io.DataInput;
import java.io.DataOutput;
import java.io.IOException;

public class CompositeKeyWritable implements WritableComparable<CompositeKeyWritable> {

    private String state;
    private String year;

    @Override
    public String toString() {
        return "CompositeKeyWritable{" +
            "state='" + state + '\'' +
            ", year='" + year + '\'' +
            '}';
    }

    public String getState() {
        return state;
    }
}
```

```
}

public void setState(String state) {
    this.state = state;
}

public String getYear() {
    return year;
}

public void setYear(String year) {
    this.year = year;
}

public CompositeKeyWritable(String state, String weatherCondition) {
    super();
    this.state = state;
    this.year = weatherCondition;
}

public CompositeKeyWritable() {
    super();
}

@Override
public int compareTo(CompositeKeyWritable o) {
    int result = this.state.compareTo(o.state);
    if(result == 0) {
        return this.year.compareTo(o.year);
    }
    return result;
}

@Override
public void write(DataOutput dataOutput) throws IOException {
    dataOutput.writeUTF(state);
    dataOutput.writeUTF(year);
}

@Override
public void readFields(DataInput dataInput) throws IOException {
```

```
        state = dataInput.readUTF();
        year = dataInput.readUTF();
    }
}
```

#### DriverClass

```
package com.mycompany.assignment_4.Q8YearandStateAnalysisTotalNumberofaccidents;
```

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
```

```
import java.io.IOException;
```

```
public class DriverClass {
```

```
    public static void main(String[] args) throws IOException, InterruptedException, ClassNotFoundException {
        Configuration config = new Configuration();
```

```
        Job job = Job.getInstance(config, "");
```

```
        job.setJarByClass(DriverClass.class);
```

```
        job.setGroupingComparatorClass(NaturalKeyComparator.class);
```

```
        job.setSortComparatorClass(CompositeKeyComparator.class);
```

```
        job.setPartitionerClass(NaturalKeyPartitioner.class);
```

```
        job.setMapOutputKeyClass(CompositeKeyWritable.class);
```

```
        job.setMapOutputValueClass(IntWritable.class);
```

```
        job.setInputFormatClass(TextInputFormat.class);
```

```
        job.setOutputFormatClass(TextOutputFormat.class);
```

```
        job.setOutputKeyClass(Text.class);
```

```
job.setOutputValueClass(IntWritable.class);

job.setMapperClass(MapperClass.class);
job.setReducerClass(ReducerClass.class);

FileInputFormat.addInputPath(job, new Path(args[0]));
Path outDir = new Path(args[1]);
FileOutputFormat.setOutputPath(job, outDir);

// Set Number of Reduce Tasks
job.setNumReduceTasks(5);

FileSystem fs = FileSystem.get(job.getConfiguration());
if (fs.exists(outDir)) {
    fs.delete(outDir, true);
}

System.exit(job.waitForCompletion(true) ? 0 : 1);
}

}
```

### MapperClass

```
package com.mycompany.assignment_4.Q8YearandStateAnalysisTotalNumberofaccidents;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

import java.io.IOException;
import java.text.ParseException;
import java.text.SimpleDateFormat;

public class MapperClass extends Mapper<Object, Text, CompositeKeyWritable, IntWritable> {

    SimpleDateFormat parser = new SimpleDateFormat("yyyy-MM-dd HH:mm:ss");
    SimpleDateFormat formater = new SimpleDateFormat("yyyy");

    @Override
    protected void map(Object key, Text value, Context context) throws IOException, InterruptedException {
```

```
String []tokens = value.toString().split(",");

String stateString = tokens[17];
String year = null;
try {
    year = formater.format(parser.parse(tokens[4]));
} catch (ParseException e) {
    e.printStackTrace();
}

CompositeKeyWritable cKW = new CompositeKeyWritable(stateString, year);

context.write(cKW, new IntWritable(1));
}
}
```

#### NaturalKeyComparator

```
package com.mycompany.assignment_4.Q8YearandStateAnalysisTotalNumberofaccidents;

import org.apache.hadoop.io.WritableComparable;
import org.apache.hadoop.io.WritableComparator;

public class NaturalKeyComparator extends WritableComparator {

    public NaturalKeyComparator() {
        super(CompositeKeyWritable.class, true);
    }

    @Override
    public int compare(WritableComparable a, WritableComparable b) {
        CompositeKeyWritable w1 = (CompositeKeyWritable) a;
        CompositeKeyWritable w2 = (CompositeKeyWritable) b;

        return w1.getState().compareTo(w2.getState());
    }
}
```

#### NaturalKeyPartitioner

```
package com.mycompany.assignment_4.Q8YearandStateAnalysisTotalNumberofaccidents;
```

```
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.mapreduce.Partitioner;

public class NaturalKeyPartitioner extends Partitioner<CompositeKeyWritable, IntWritable> {

    @Override
    public int getPartition(CompositeKeyWritable compositeKeyWritable, IntWritable intWritable, int i) {
        return Integer.parseInt(compositeKeyWritable.getYear()) % i;
    }
}

ReducerClass

package com.mycompany.assignment_4.Q8YearandStateAnalysisTotalNumberofaccidents;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

import java.io.IOException;

public class ReducerClass extends Reducer<CompositeKeyWritable, IntWritable, Text, IntWritable> {

    Text text = new Text();
    IntWritable count = new IntWritable();
    StateMap stateMap = new StateMap();

    @Override
    protected void reduce(CompositeKeyWritable key, Iterable<IntWritable> values, Context context) throws
    IOException, InterruptedException {

        int sum = 0;

        for(IntWritable v : values) {
            sum += v.get();
        }

        text.set(stateMap.getStateFromAbbr(key.getState())+ " - in Year - " + key.getYear() + " : ");
        count.set(sum);

        context.write(text, count);
    }
}
```

```
}  
}
```

### StateMap

```
package com.mycompany.assignment_4.Q8YearandStateAnalysisTotalNumberofaccidents;
```

```
import java.util.HashMap;
```

```
import java.util.Map;
```

```
public class StateMap {
```

```
    public static final Map<String, String> STATE_MAP;
```

```
    static {
```

```
        STATE_MAP = new HashMap<String, String>();
```

```
        STATE_MAP.put("AL", "Alabama");
```

```
        STATE_MAP.put("AK", "Alaska");
```

```
        STATE_MAP.put("AB", "Alberta");
```

```
        STATE_MAP.put("AZ", "Arizona");
```

```
        STATE_MAP.put("AR", "Arkansas");
```

```
        STATE_MAP.put("BC", "British Columbia");
```

```
        STATE_MAP.put("CA", "California");
```

```
        STATE_MAP.put("CO", "Colorado");
```

```
        STATE_MAP.put("CT", "Connecticut");
```

```
        STATE_MAP.put("DE", "Delaware");
```

```
        STATE_MAP.put("DC", "District Of Columbia");
```

```
        STATE_MAP.put("FL", "Florida");
```

```
        STATE_MAP.put("GA", "Georgia");
```

```
        STATE_MAP.put("GU", "Guam");
```

```
        STATE_MAP.put("HI", "Hawaii");
```

```
        STATE_MAP.put("ID", "Idaho");
```

```
        STATE_MAP.put("IL", "Illinois");
```

```
        STATE_MAP.put("IN", "Indiana");
```

```
        STATE_MAP.put("IA", "Iowa");
```

```
        STATE_MAP.put("KS", "Kansas");
```

```
        STATE_MAP.put("KY", "Kentucky");
```

```
        STATE_MAP.put("LA", "Louisiana");
```

```
        STATE_MAP.put("ME", "Maine");
```

```
        STATE_MAP.put("MB", "Manitoba");
```

```
        STATE_MAP.put("MD", "Maryland");
```

```
        STATE_MAP.put("MA", "Massachusetts");
```

```
STATE_MAP.put("MI", "Michigan");
STATE_MAP.put("MN", "Minnesota");
STATE_MAP.put("MS", "Mississippi");
STATE_MAP.put("MO", "Missouri");
STATE_MAP.put("MT", "Montana");
STATE_MAP.put("NE", "Nebraska");
STATE_MAP.put("NV", "Nevada");
STATE_MAP.put("NB", "New Brunswick");
STATE_MAP.put("NH", "New Hampshire");
STATE_MAP.put("NJ", "New Jersey");
STATE_MAP.put("NM", "New Mexico");
STATE_MAP.put("NY", "New York");
STATE_MAP.put("NF", "Newfoundland");
STATE_MAP.put("NC", "North Carolina");
STATE_MAP.put("ND", "North Dakota");
STATE_MAP.put("NT", "Northwest Territories");
STATE_MAP.put("NS", "Nova Scotia");
STATE_MAP.put("NU", "Nunavut");
STATE_MAP.put("OH", "Ohio");
STATE_MAP.put("OK", "Oklahoma");
STATE_MAP.put("ON", "Ontario");
STATE_MAP.put("OR", "Oregon");
STATE_MAP.put("PA", "Pennsylvania");
STATE_MAP.put("PE", "Prince Edward Island");
STATE_MAP.put("PR", "Puerto Rico");
STATE_MAP.put("QC", "Quebec");
STATE_MAP.put("RI", "Rhode Island");
STATE_MAP.put("SK", "Saskatchewan");
STATE_MAP.put("SC", "South Carolina");
STATE_MAP.put("SD", "South Dakota");
STATE_MAP.put("TN", "Tennessee");
STATE_MAP.put("TX", "Texas");
STATE_MAP.put("UT", "Utah");
STATE_MAP.put("VT", "Vermont");
STATE_MAP.put("VI", "Virgin Islands");
STATE_MAP.put("VA", "Virginia");
STATE_MAP.put("WA", "Washington");
STATE_MAP.put("WV", "West Virginia");
STATE_MAP.put("WI", "Wisconsin");
STATE_MAP.put("WY", "Wyoming");
STATE_MAP.put("YT", "Yukon Territory");
```



```

    }

    public String getStateFromAbbr(String abbr) {
        return STATE_MAP.get(abbr);
    }

}

```

## Q9. Percentage Accidents vs Proximity

This MapReduce Numerical Summarization performs the percentage of accidents vs. the proximity to traffic object.

```

nidhi@nidhi: /usr/local/bin/hadoop-2.7.3/bin$ ./hadoop jar /home/nidhi/Desktop/satwik_bigdata/Q9.jar com.hadoop.FinalProject.Q11.DriverClass /Info7250_2020 /SatwikResults/Q9
20/12/17 21:33:20 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
20/12/17 21:33:20 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
20/12/17 21:33:21 INFO Input.FileInputFormat: Total input paths to process : 1
20/12/17 21:33:21 INFO mapreduce.JobSubmitter: number of splits:10
20/12/17 21:33:21 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1608248153408_0009
20/12/17 21:33:21 INFO Impl.YarnClientImpl: Submitted application application_1608248153408_0009
20/12/17 21:33:21 INFO mapreduce.Job: The url to track the job: https://nidhi:8088/proxy/application_1608248153408_0009/
20/12/17 21:33:21 INFO mapreduce.Job: Running job: job_1608248153408_0009
20/12/17 21:33:31 INFO mapreduce.Job: Job job_1608248153408_0009 running in uber mode : false
20/12/17 21:33:31 INFO mapreduce.Job:  map 0% reduce 0%
20/12/17 21:33:57 INFO mapreduce.Job:  map 10% reduce 0%
20/12/17 21:33:58 INFO mapreduce.Job:  map 13% reduce 0%
20/12/17 21:34:00 INFO mapreduce.Job:  map 28% reduce 0%
20/12/17 21:34:01 INFO mapreduce.Job:  map 31% reduce 0%
20/12/17 21:34:03 INFO mapreduce.Job:  map 39% reduce 0%
20/12/17 21:34:04 INFO mapreduce.Job:  map 53% reduce 0%
20/12/17 21:34:05 INFO mapreduce.Job:  map 68% reduce 0%
20/12/17 21:34:21 INFO mapreduce.Job:  map 75% reduce 0%
20/12/17 21:34:22 INFO mapreduce.Job:  map 80% reduce 0%
20/12/17 21:34:23 INFO mapreduce.Job:  map 85% reduce 23%
20/12/17 21:34:24 INFO mapreduce.Job:  map 100% reduce 23%
20/12/17 21:34:26 INFO mapreduce.Job:  map 100% reduce 82%
20/12/17 21:34:27 INFO mapreduce.Job:  map 100% reduce 100%
20/12/17 21:34:27 INFO mapreduce.Job: Job job_1608248153408_0009 completed successfully
20/12/17 21:34:27 INFO mapreduce.Job: Counters: 50
  File System Counters
    FILE: Number of bytes read=24307994
    FILE: Number of bytes written=49923293
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=1327235602
    HDFS: Number of bytes written=196
    HDFS: Number of read operations=33
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Killed map tasks=2
    Launched map tasks=11
    Launched reduce tasks=1
    Data-local map tasks=11
    Total time spent by all maps in occupied slots (ms)=264253
    Total time spent by all reduces in occupied slots (ms)=20885
    Total time spent by all map tasks (ms)=264253
    Total time spent by all reduce tasks (ms)=20885
    Total vcore-millisecons taken by all map tasks=264253
    Total vcore-millisecons taken by all reduce tasks=20885
    Total megabyte-millisecons taken by all map tasks=270595072
    Total megabyte-millisecons taken by all reduce tasks=21386240
  Map-Reduce Framework
    Map input records=3513618
    Map output records=1394291
    Map output bytes=21519406
    Map output materialized bytes=24308048
    Input split bytes=1000
    Combine input records=0
    Combine output records=0
    Reduce input groups=12
    Reduce shuffle bytes=24308048
    Reduce input records=1394291
    Reduce output records=12
    Spilled Records=2788582
    Shuffled Maps =10
    Failed Shuffles=0
    Merged Map outputs=10
    GC time elapsed (ms)=9590
    CPU time spent (ms)=76670
    Physical memory (bytes) snapshot=2934112256
    Virtual memory (bytes) snapshot=20610441216
    Total committed heap usage (bytes)=2184183808
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=1327234602
  File Output Format Counters
    Bytes Written=196
nidhi@nidhi: /usr/local/bin/hadoop-2.7.3/bin$ ./hadoop fs -tail /SatwikResults/Q9/part-r-00000 | head
hadoopfs 1.20%
hadoopfs 0.02%

```

```

nidhi@nidhi:/usr/local/bin/hadoop-2.7.3/bin$ ./hadoop fs -tail /SatwikResults/Q9/part-r-00000 | head
Amenity          1.20%
Bump              0.02%
Crossing          7.81%
Give_Way          0.27%
Junction          8.10%
No_Exit           0.12%
Railway           0.89%
Roundabout        0.01%
Station           2.00%
Stop              1.48%
nidhi@nidhi:/usr/local/bin/hadoop-2.7.3/bin$

```

## Driver

```
package com.mycompany.assignment_4.Q9PercentageAccidentsvsProximity;
```

```

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

```

```
import java.io.IOException;
```

```
public class DriverClass {
```

```

    public static void main(String[] args) throws IOException, InterruptedException, ClassNotFoundException {
        Configuration conf = new Configuration();

```

```
        Job job = Job.getInstance(conf, "ProximityToTrafficObject");
```

```
        FileInputFormat.addInputPath(job, new Path(args[0]));
```

```
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

```
        // Driver Class
```

```
        job.setJarByClass(DriverClass.class);
```

```
        job.setInputFormatClass(TextInputFormat.class);
```

```
        Path outDir = new Path(args[1]);
```

```
        FileOutputFormat.setOutputPath(job, outDir);
```

```
// Mapper
job.setMapperClass(MapperClass.class);

job.setMapOutputKeyClass(Text.class);
job.setMapOutputValueClass(IntWritable.class);

// Reducer
job.setReducerClass(ReducerClass.class);

// Create only 1 reducer
job.setNumReduceTasks(1);

job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);

FileSystem fs = FileSystem.get(job.getConfiguration());
if(fs.exists(outDir)){
    fs.delete(outDir, true);
}

// Submit the job, then poll for progress until the job is complete
System.exit(job.waitForCompletion(true) ? 0 : 1);
}

}
```

## Mapper

```
package com.mycompany.assignment_4.Q9PercentageAccidentsvsProximity;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

import java.io.IOException;

public class MapperClass extends Mapper<LongWritable, Text, Text, IntWritable> {

    Text proximity = new Text();
```

```
IntWritable one = new IntWritable(1);
```

```
RowHeaderClass rhc = new RowHeaderClass();
```

```
@Override
```

```
protected void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
```

```
    String[] tokens = value.toString().split(",");
```

```
    if (!tokens[0].equals("ID")) {
```

```
        for(int i = 32; i < 45; i++) {
```

```
            if(tokens[i].equals("True")) {
```

```
                proximity.set(rhc.getRowName(i));
```

```
                context.write(proximity, one);
```

```
            }
```

```
        }
```

```
    }
```

```
}
```

```
}
```

Reducer

```
package com.mycompany.assignment_4.Q9PercentageAccidentsvsProximity;
```

```
import org.apache.hadoop.io.IntWritable;
```

```
import org.apache.hadoop.io.Text;
```

```
import org.apache.hadoop.mapreduce.Reducer;
```

```
import java.io.IOException;
```

```
public class ReducerClass extends Reducer<Text, IntWritable, Text, Text> {
```

```
    Text percentage = new Text();
```

```
@Override
```

```
protected void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException,
InterruptedException {
```

```
    int count = 0;
```

```
    for(IntWritable v : values) {
```

```
        count += v.get();
```

```
    }

    int total = 3513617; // total number of records

    double perc = ((double) count / total) * 100;
    percentage.set("\t" + String.format("%.2f", perc) + "%");

    context.write(key, percentage);
  }
}
```

#### RowHeaders

```
package com.mycompany.assignment_4.Q9PercentageAccidentsvsProximity;

import java.util.HashMap;
import java.util.Map;

public class RowHeaderClass {

    Map<Integer, String> rowNames = new HashMap<>();
    String []headerNames =
{"Amenity","Bump","Crossing","Give_Way","Junction","No_Exit","Railway","Roundabout","Station","Stop","Traffic_Cal
ming","Traffic_Signal","Turning_Loop"};

    public RowHeaderClass() {

        for(int i = 32; i < 45; i++) {
            rowNames.put(i, headerNames[i-32]);
        }

    }

    public String getRowName(int rowNumber) {
        return rowNames.get(rowNumber);
    }
}
```

**Q1. Number of Accidents per Time Zone**

---

**Q2. Weather Conditions leading to accidents**

---

**References**

---

<https://www.kaggle.com/sobhanmoosavi/us-accidents>

<https://stackoverflow.com/>

<https://american-cse.org/csci2015/data/9795a392.pdf>

<https://towardsdatascience.com/usa-accidents-data-analysis-d130843cde02>